
Introductory Session: Applied Deep Learning

August 17, 2017

Dr. John See, Multimedia University (MMU)

Dr. Poo Kuan Hoong, ADAX



Agenda

- **Introduction**
 - Machine Learning
- **Deep Learning**
 - Artifical Neural Network
 - Convolution Neural Network (CNN)
 - Recurrent Neural Network (RNN)
 - Generative Adversarial Network (GAN)
 - Use Cases
 - Deep Learning Libraries/Tools
- **Demo - TensorFlow & Keras**
- **Course Preview**

About “Applied Deep Learning”

This is an **introductory preview** towards a **course** aimed at

- Graduate students,
- Researchers,
- IT professionals
- Industrial engineers

who already possess **basic knowledge in Python programming and machine learning** (and possibly, but not necessarily of deep learning) - who wish to learn more about this rapidly growing field of research and its application in the commercial world.

About us:



John See, Ph.D <http://pesona.mmu.edu.my/~johnsee>



Senior Lecturer, Faculty of Computing and
Informatics
Head, Visual Processing Lab
(<http://viprlab.github.io>)
Theme Lead, Smart Cities & Digital Services

Co-Organizer,
TensorFlow and Deep Learning Malaysia User
Group



About us:



Poo Kuan Hoong, Ph.D

<http://www.linkedin.com/in/kuanhoong>



- Senior Data Scientist



- Senior Manager Data Science



- Senior Lecturer
- Chairperson Data Science Institute



- Founder R User Group & TensorFlow User Group
- Speaker/Trainer

Where are we now..

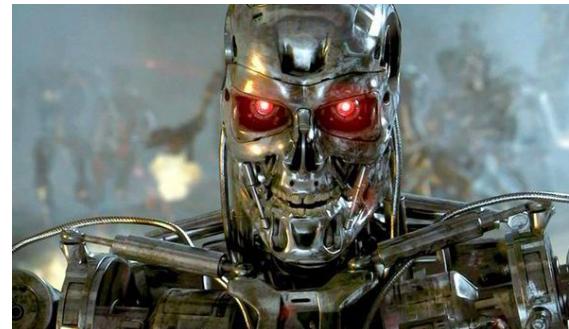
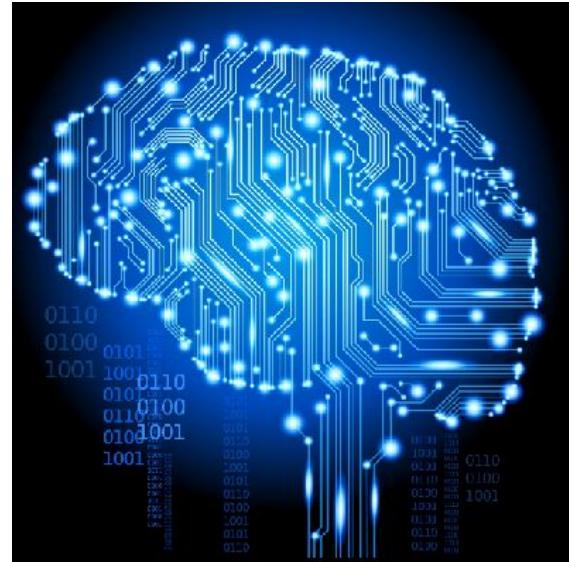


Where are we now..



Introduction

- **Deep Learning:** The fastest growing sub-field in A.I.
 - Computers can now make sense of **difficult tasks** by crunching **large amounts of data**
- **See, Hear, Learn, React**
- Many industries and businesses have been impacted by DL, spurring new products and services that are **data-driven**
- **Machine Learning:** Machines that can learn!

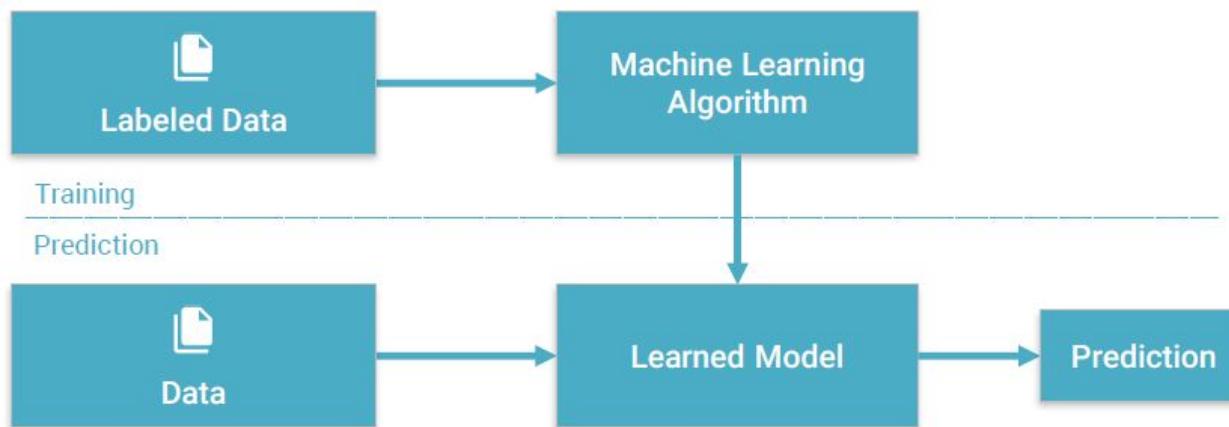




How do we learn?

Machine Learning

A type of Artificial Intelligence that provides computers with the ability to learn **without being explicitly programmed**



Machine Learning - Approaches



Supervised Learning

Learning from a labeled training set



Unsupervised Learning

Discovering patterns in unlabeled data



Reinforcement Learning

Learning based on feedback or reward

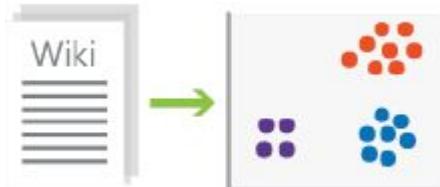
Machine Learning - Types of Problems



Classification



Regression



Clustering



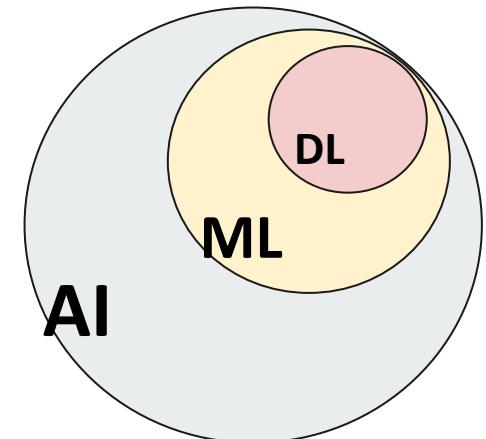
Anomaly Detection

Deep Learning



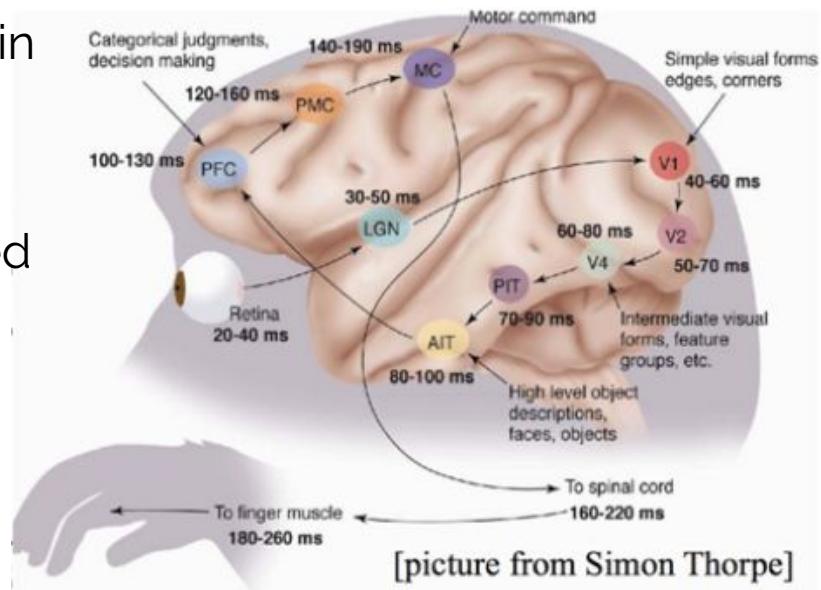
Deep Learning

- In the past 10 years, Artificial Intelligence (AI) particularly ML have shown tremendous progress.
- **Deep Learning**: part of Machine Learning (ML) field of **learning representations of data**
 - Exceptional effectiveness
 - Hierarchy of multiple “learning” layers that mimic the neurons in our brain
 - Thrives on big amounts of data



Why learn “deeply”?

- Inspired by architectural depth of brain
- Neuroscience findings: Information becomes more complex down the visual pipeline
- No successful attempts were reported before 2006
 - SVM: Vapnik and co. developed Support Vector Machine which was pretty good, but shallow learning

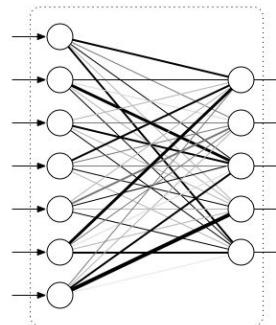


2006 Breakthrough: What made Deep Learning possible?

- Explosion of data
- Faster and cheaper computing -- multi-core CPUs and GPUs
- Improvement of ML models -- more complex, more scalable



Stacked Restricted Boltzmann Machines (RBM)
or Deep Belief Nets (DBN), Hinton, 2006
Stacked Autoencoders (AE), Bengio, 2007



“I've worked all my life in Machine Learning, and I've never seen one algorithm knock over benchmarks like Deep Learning”

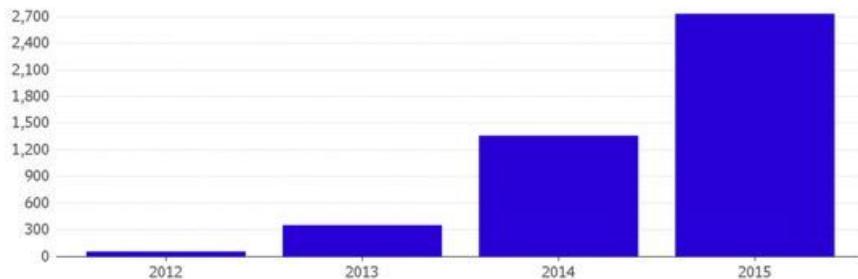
Andrew Ng
Stanford Univ, Baidu (prev.)
Co-Founder of Coursera



Deep Learning: Hype or Reality?

Artificial Intelligence Takes Off at Google

Number of software projects within Google that uses a key AI technology, called Deep Learning.



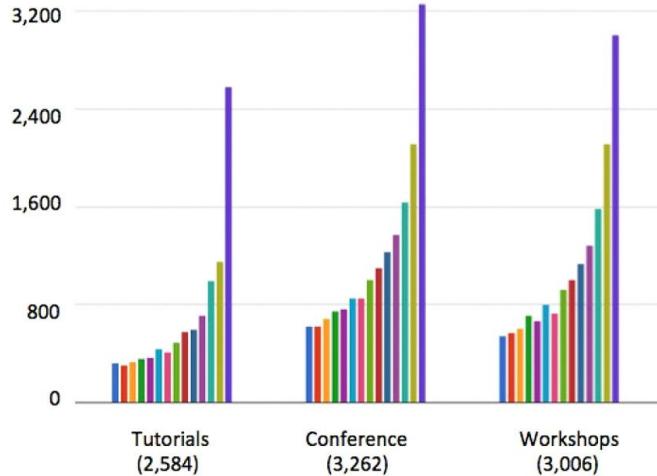
Source: Google

Note: 2015 data does not incorporate data from Q4

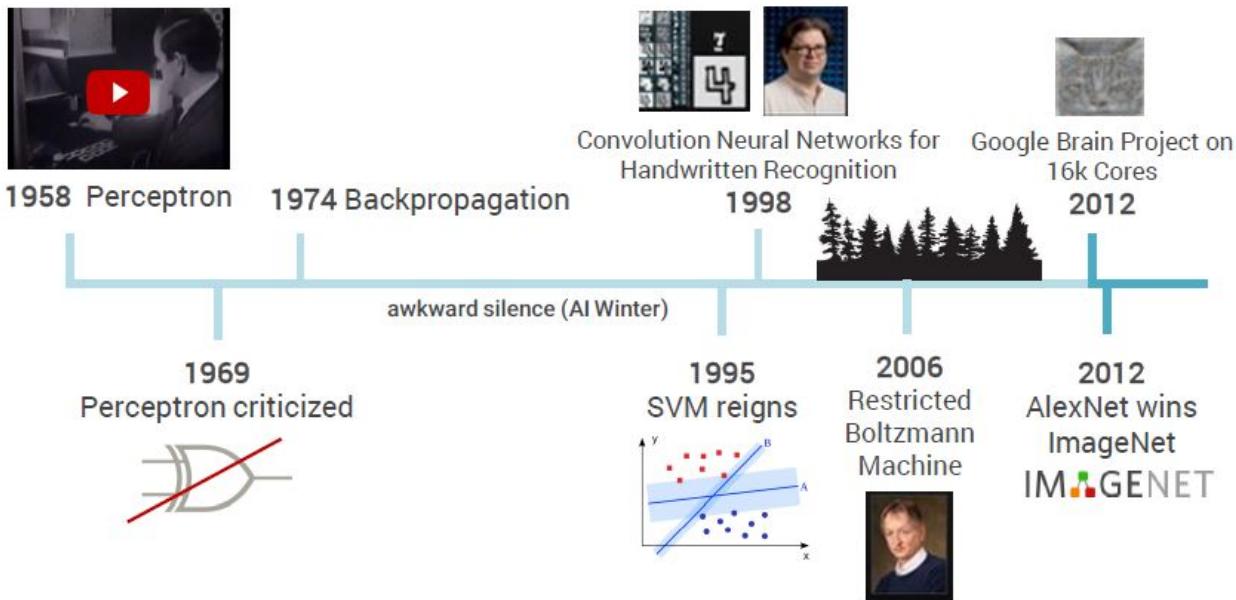
Bloomberg

NIPS Growth

Total Registrations 3755

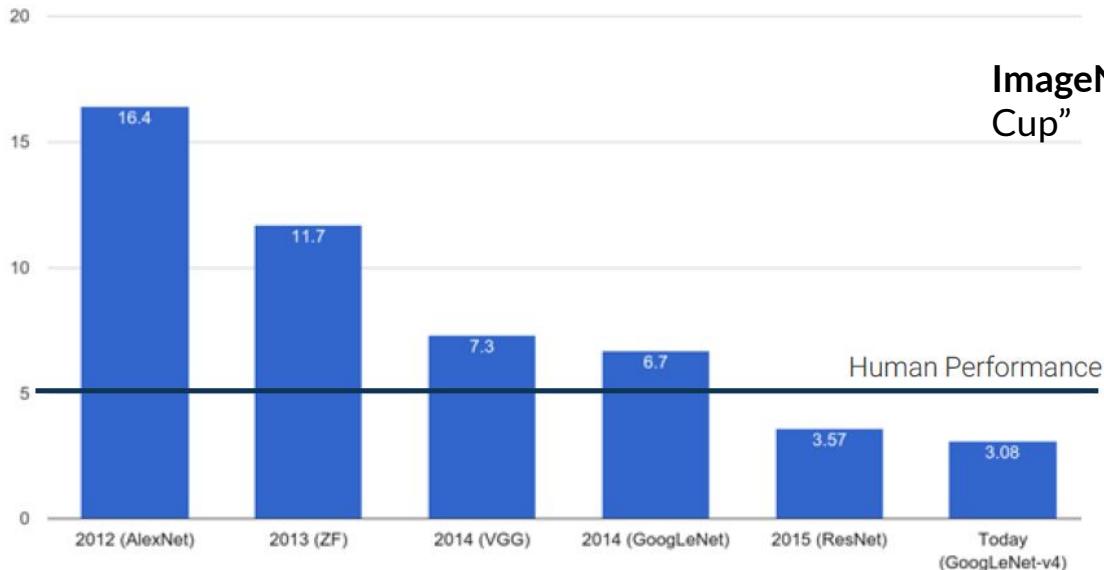


A very brief history...



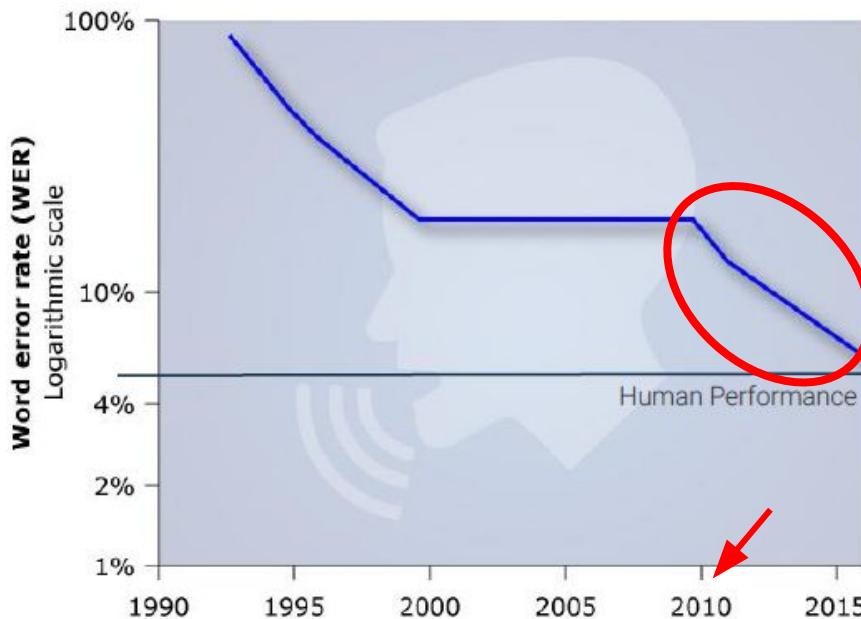
“One algorithm to rule them all”

ImageNet Classification Error (Top 5)



ImageNet: The “computer vision World Cup”

“One algorithm to rule them all”

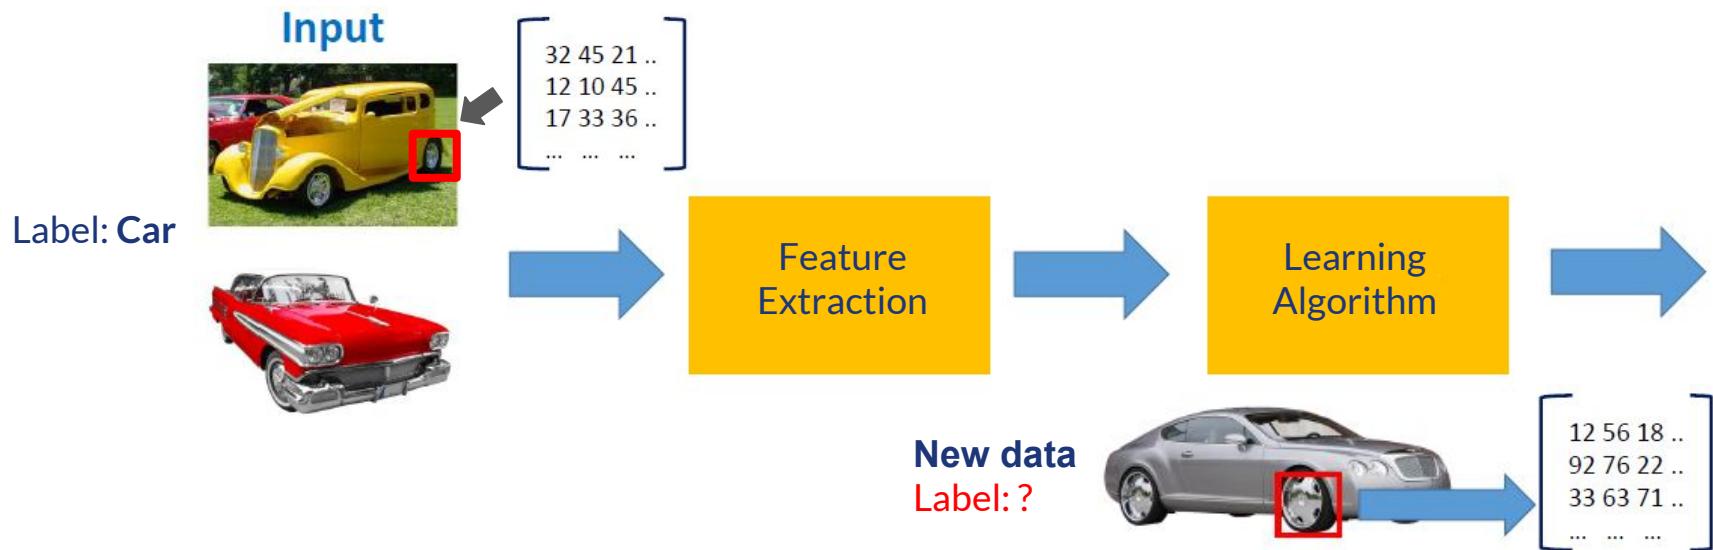


State-of-the-art Performance for
Speech Recognition

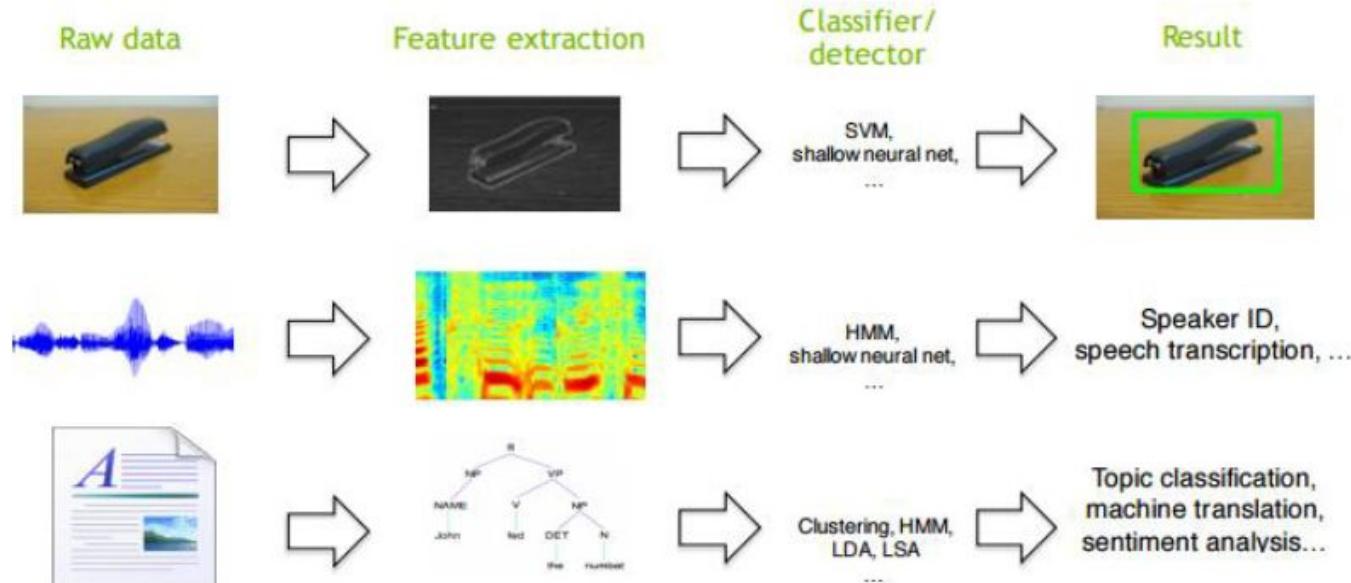
Thanks
to
DL!

Revisiting Traditional Machine Learning

- **“Feature Engineering”:** Relates to extracting features by “hand-crafted” means. A simple classifier can be trained for recognition.

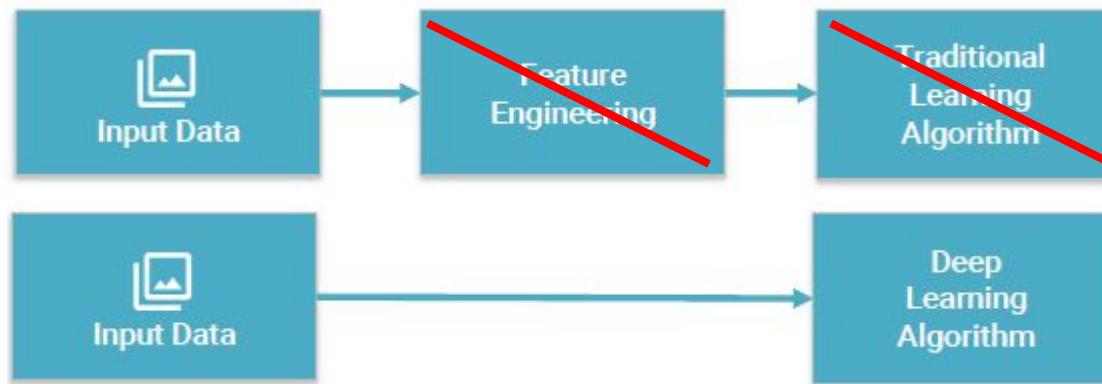


Revisiting Traditional Machine Learning



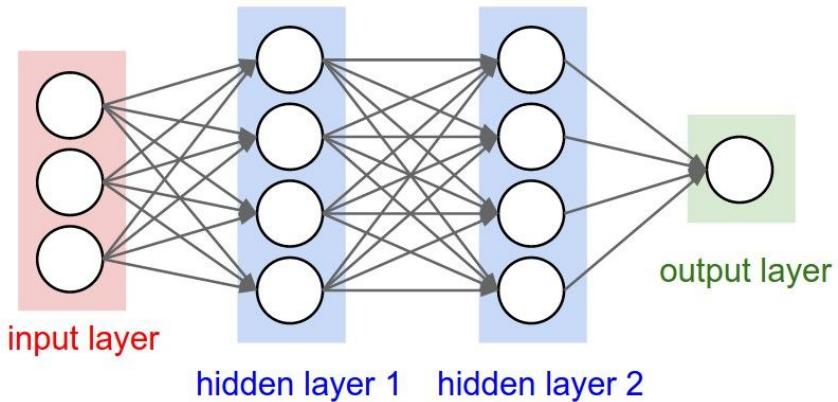
Deep Learning - All-in-one

- Removes the need to **manually craft** the feature (DL “**finds**” them instead)
- Removes the need to **perform separate learning** (DL “**learns**” them directly)

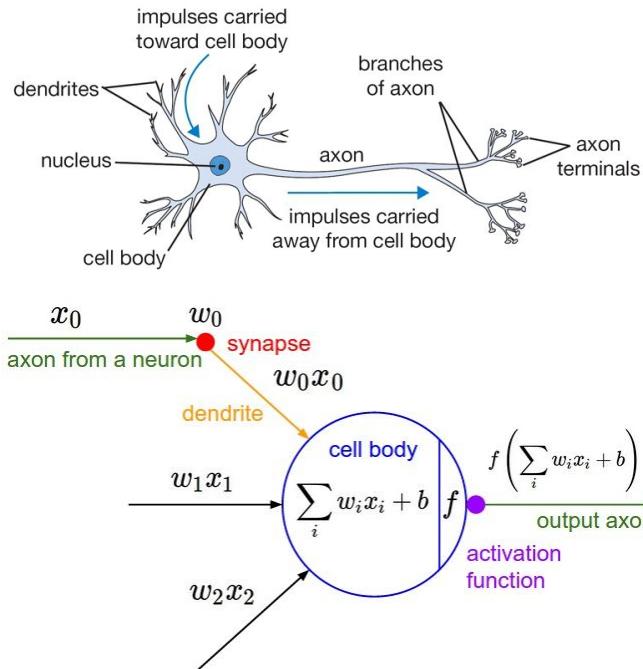


Artificial Neural Networks

- An **artificial neural network** consists of **one input, one output** and **multiple fully-connected hidden layers** in between
- Each layer -- a series of neurons, which connects the previous layer to the next

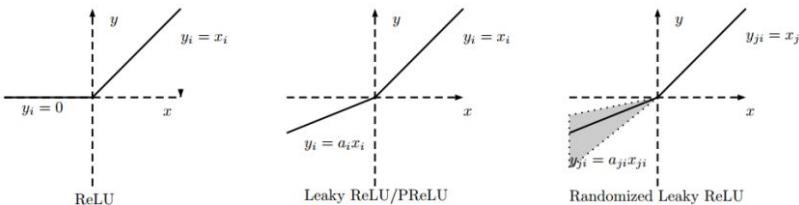
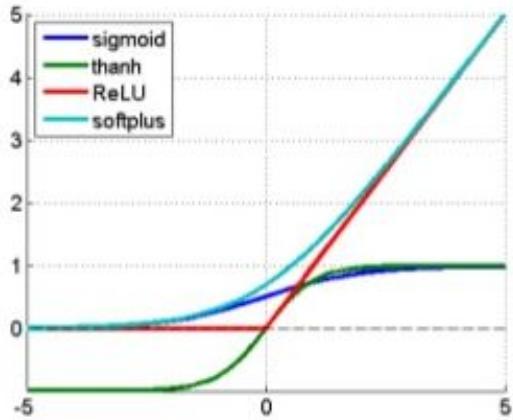


NN Basics: The Neuron



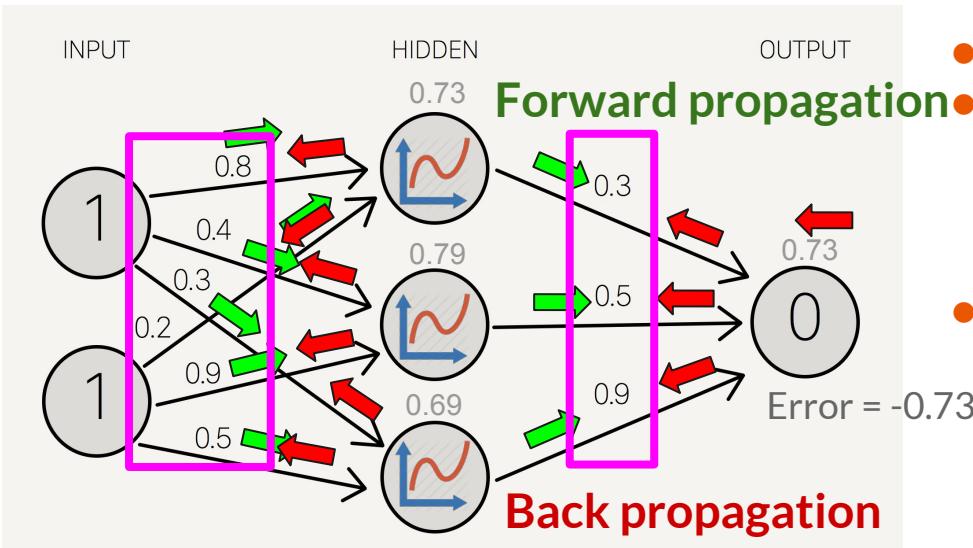
- An **artificial neuron** contains a **nonlinear activation function** and has several incoming and outgoing **weighted connections**
- **Neurons**: trained to filter and detect specific features or patterns

NN Basics: Activation Functions



- Why do activation functions need to be **non-linear**?
 - To learn complex representations of data
- Among the more popular choices in DL: **ReLU (Rectified Linear Units)**
 - Creates sparsity
 - Prevents vanishing gradient problem

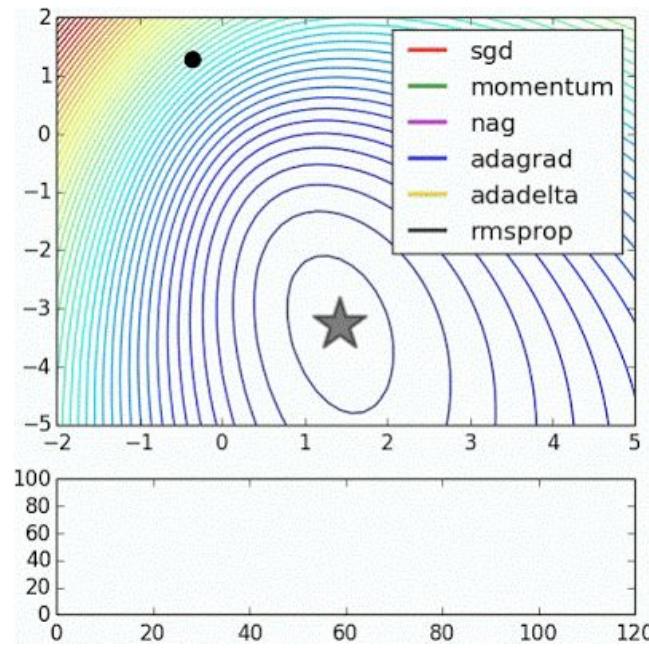
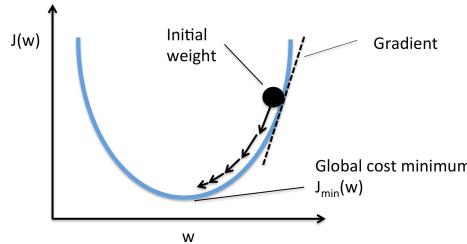
NN Basics: Forward- and Back-propagation



- Multilayer Perceptron (MLP)
 - Forward propagation
 - Sum inputs, produce activations at each step forward
 - Back propagation
 - Calculates total error, and contributions to error at each step backwards

NN Basics: Optimizing the Cost Function

- Re-adjusting the weights and biases is an **optimization** problem
- A **cost function** is defined to measure overall error which needs to be minimized
- **Gradients** are used to update weight parameters



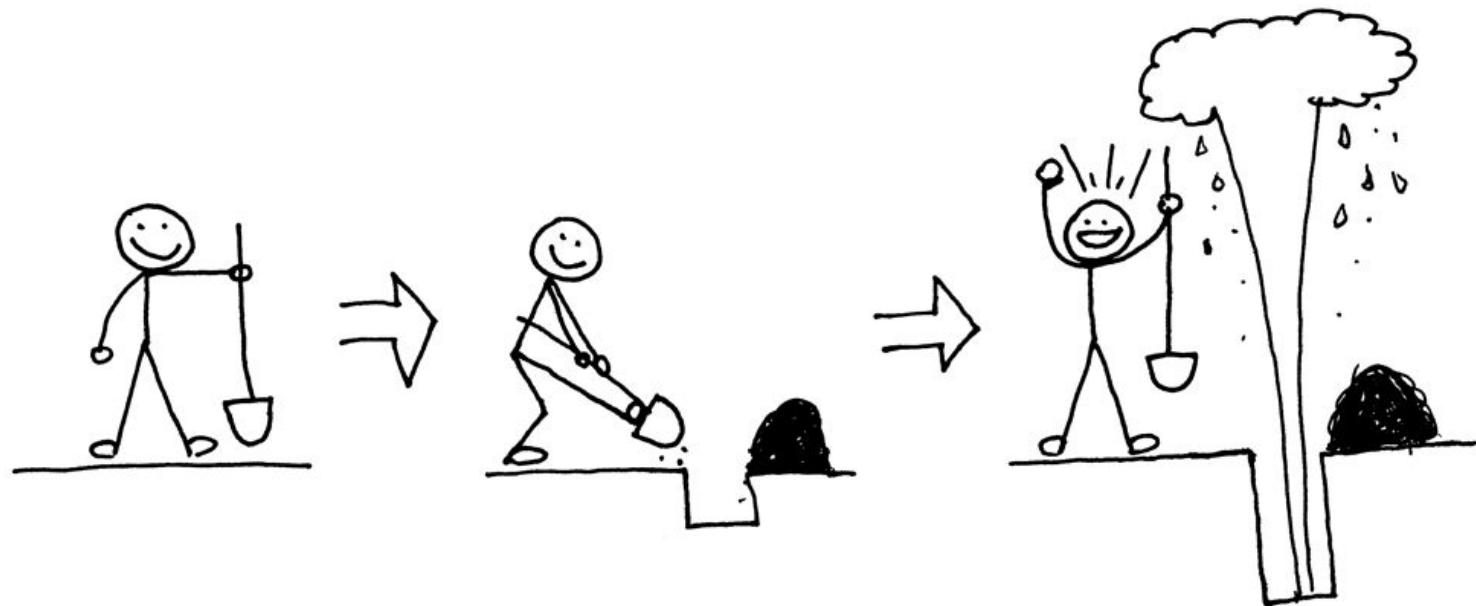
What's wrong with NN and back-propagation?

- It requires a lot of labelled training data
 - Almost all data is unlabelled
- The learning time does not scale well
 - It is very slow in networks with multiple hidden layers
- It can get stuck in poor local optimas
 - Works well most of the time, but suffers when complexity grows

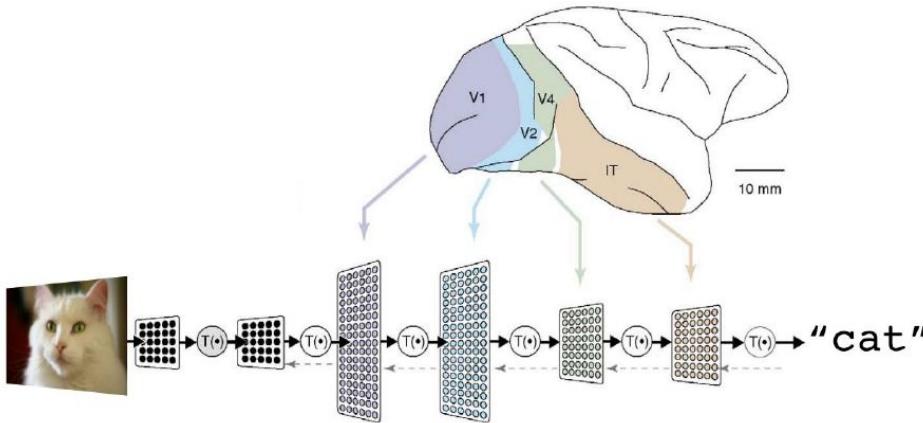
Use unsupervised
generative methods

Bring out the GPUs!

Use more robust
activation functions
like ReLU, dropout
layers, etc.



Why Go Deeper in Neural Networks?

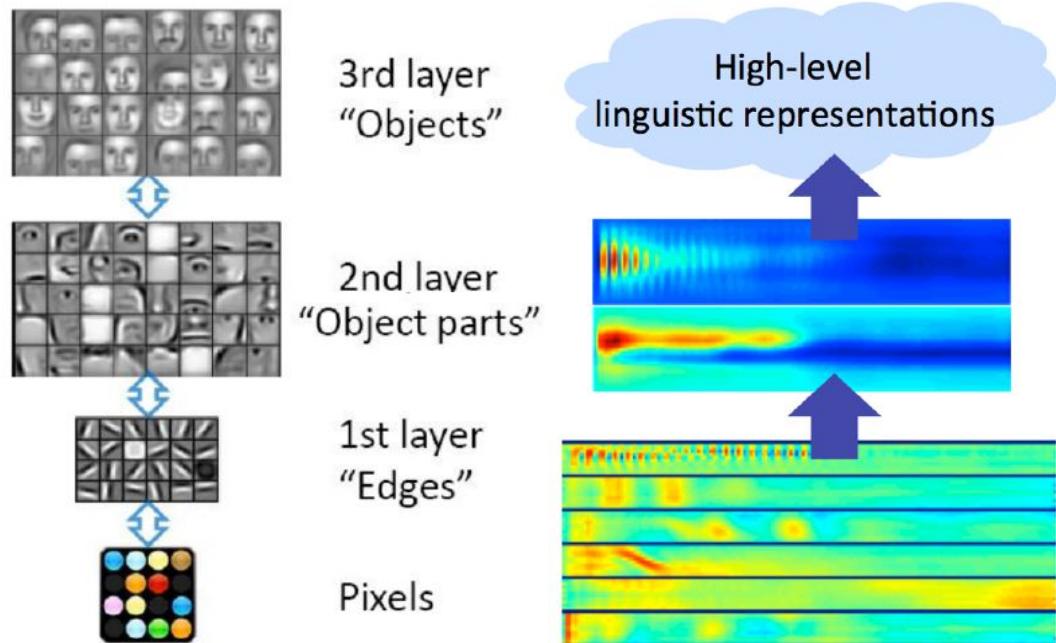


1. Biological Justification -- Hierarchical learning

- A deep neural network consists of a hierarchy of layers, whereby each layer transforms input data into more abstract representations. The output layer combines these features to make predictions

Why Go Deeper in Neural Networks?

2. Different levels of abstraction
 - Natural progression from low to high level structure as seen in natural complexity

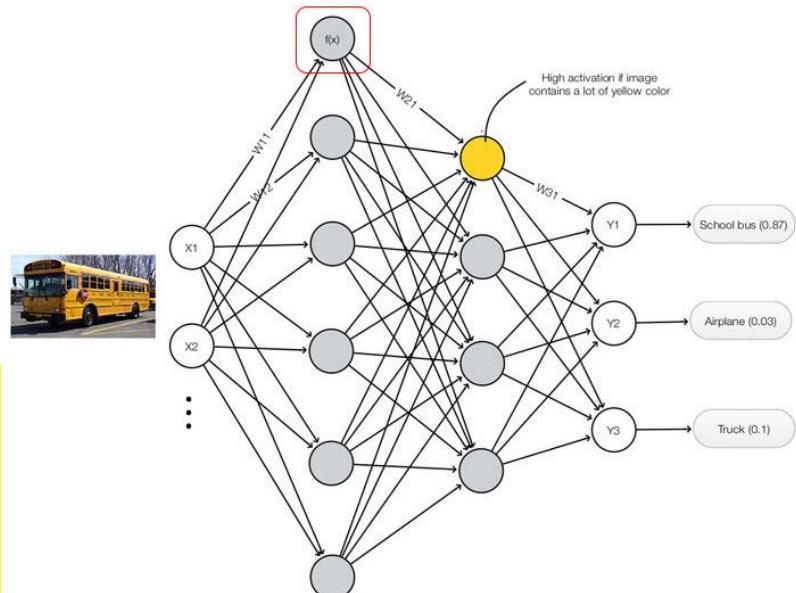


How “deep” is a deep neural network?

- Credit Assignment Path (**CAP**)
or (number of hidden layers +
1) **more than 2**
- CAP > 10 : Very deep



“There is no universally agreed upon threshold of depth dividing shallow learning from deep learning, but most researchers in the field agree that deep learning has multiple nonlinear layers ($CAP > 2$) and Schmidhuber considers $CAP > 10$ to be very deep learning”



Deep Learning: Strengths

- **Robust**
 - Features are automatically learned to be optimal for the task at hand
- **Generalizable**
 - The same DL approach/architecture can be used for many different applications and data types
- **Scalable**
 - Performance improves with more data, method is massively parallelizable

Deep Learning: Weaknesses

- Requires large datasets (typically), also means long training periods.
- Not necessarily > traditional ML methods like SVMs, decision trees
- Learned features are often difficult to understand or explain.
 - Many vision features are also not really human-understandable (e.g, concatenations/combinations of different features).
- Requires a good understanding of how to design new models for new tasks, or using new forms of data, etc.

DL Flavours



Markov Chain (MC)



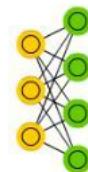
Hopfield Network (HN)



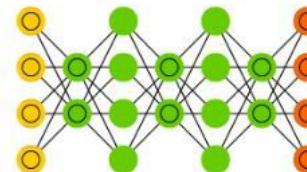
Boltzmann Machine (BM)



Restricted BM (RBM)

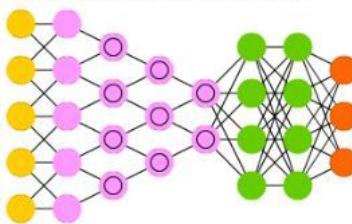


Deep Belief Network (DBN)

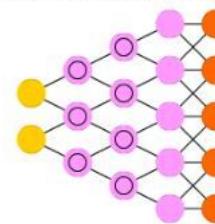


- Backfed Input Cell
- Input Cell
- △ Noisy Input Cell
- Hidden Cell
- Probabilistic Hidden Cell
- △ Spiking Hidden Cell
- Output Cell
- Match Input Output Cell
- Recurrent Cell
- Memory Cell
- △ Open Memory Cell
- Scanning Filter
- Convolution

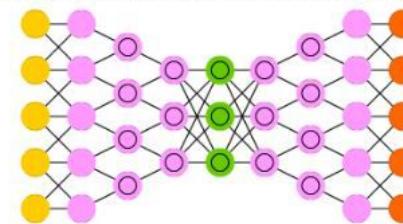
Deep Convolutional Network (DCN)



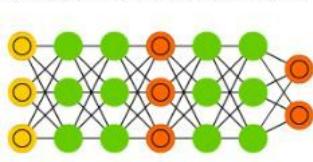
Deconvolutional Network (DN)



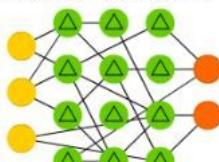
Deep Convolutional Inverse Graphics Network (DCIGN)



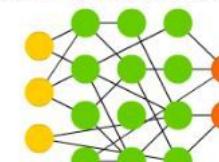
Generative Adversarial Network (GAN)



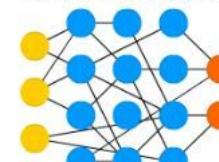
Liquid State Machine (LSM)



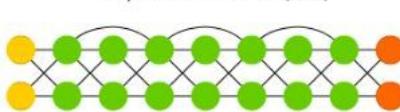
Extreme Learning Machine (ELM)



Echo State Network (ESN)



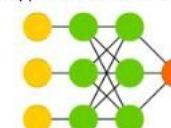
Deep Residual Network (DRN)



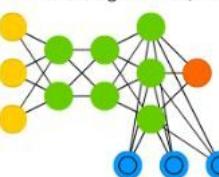
Kohonen Network (KN)



Support Vector Machine (SVM)

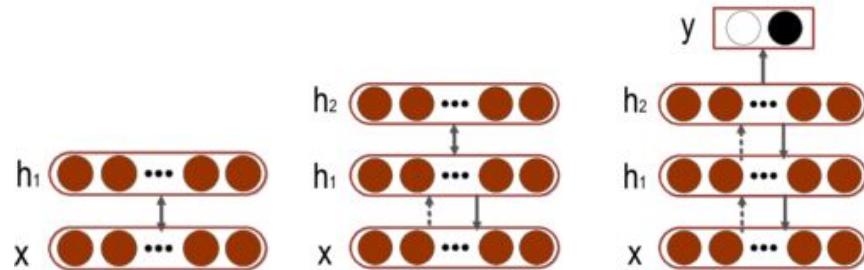


Neural Turing Machine (NTM)



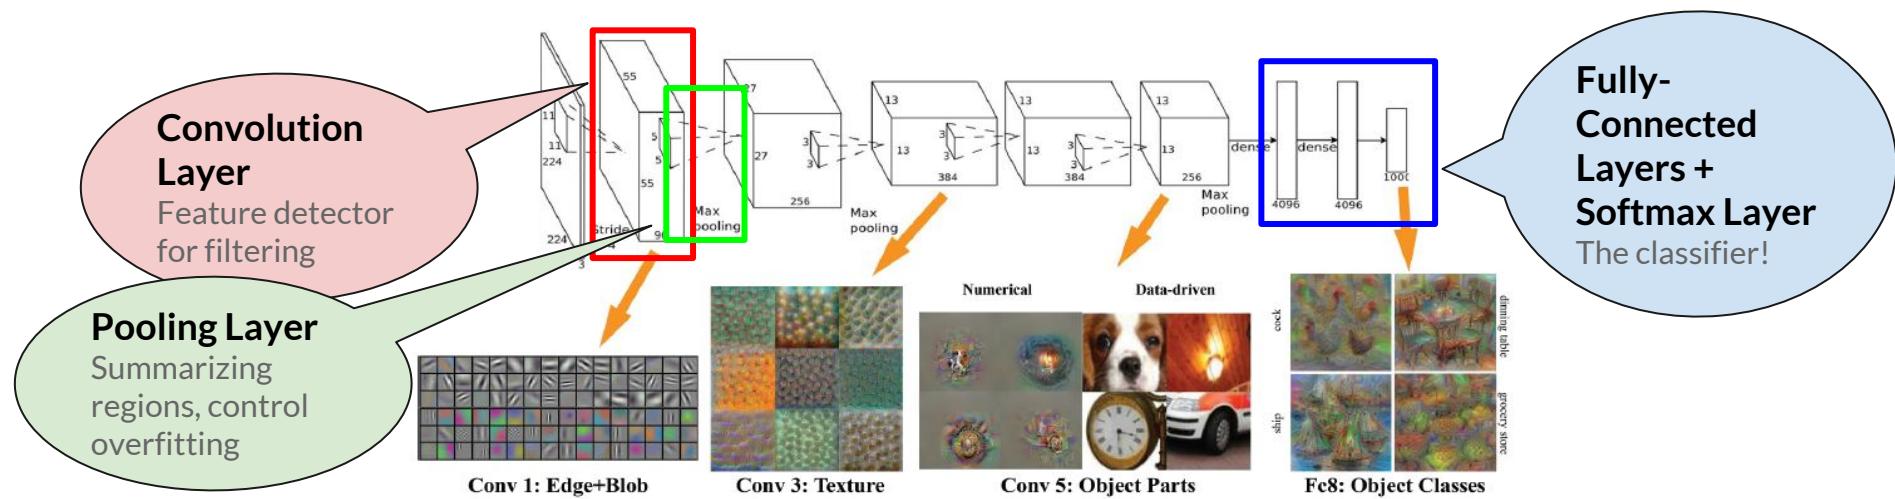
Stacking Neural Networks

- Early attempts, pioneered by Geoffrey Hinton:
 - Deep Belief Networks (DBN) or Stacked Restricted Boltzmann Machines (RBM),
 - Stacked AutoEncoders (SAE)
- Idea:
 - Layer-wise Unsupervised Pre-training
 - Train last layer in supervised mode



Convolutional Neural Networks (CNN)

- Learn multiple layers of transformations, stacked to extract a progressively more sophisticated representation of the input



Convolutional Neural Networks (CNN)

- Each neuron in the Convolutional layer is mapped to a receptive field

0	0	0	0	0	0	0	...
0	156	155	156	158	158	...	
0	153	154	157	159	159	...	
0	149	151	155	158	159	...	
0	146	146	149	153	158	...	
0	145	143	143	148	158	...	
...	

Input Channel #1 (Red)

0	0	0	0	0	0	0	...
0	167	166	167	169	169	...	
0	164	165	168	170	170	...	
0	160	162	166	169	170	...	
0	156	156	159	163	168	...	
0	155	153	153	158	168	...	
0	154	152	152	157	167	...	
...	

Input Channel #2 (Green)

0	0	0	0	0	0	0	...
0	163	162	163	165	165	...	
0	160	161	164	166	166	...	
0	156	158	162	165	166	...	
0	155	155	158	162	167	...	
0	154	152	152	157	167	...	
0	153	153	153	158	168	...	
...	

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1

308

+

-498

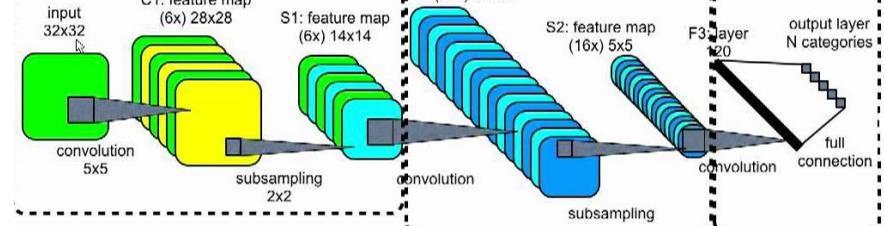
+

164

Bias = 1

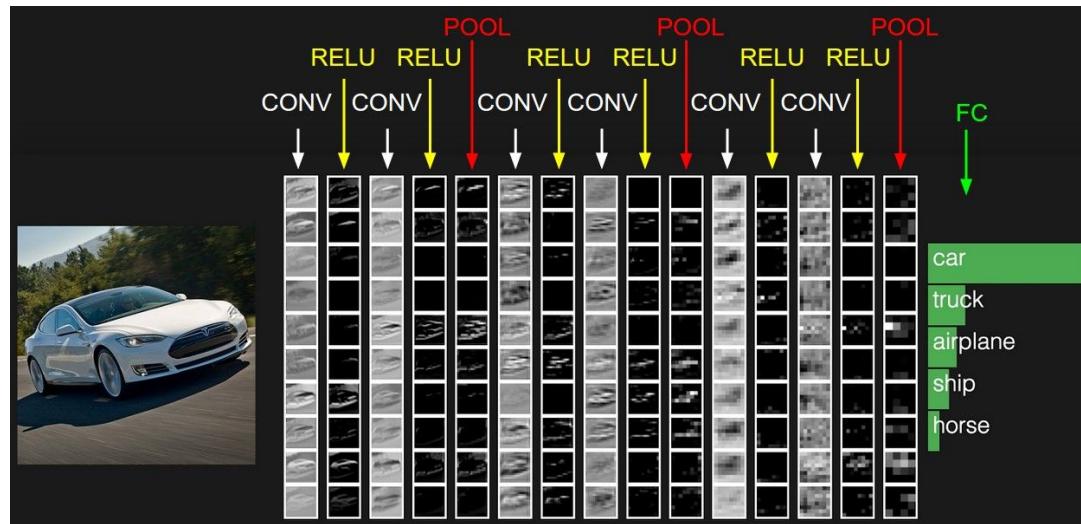
-25			...
			...
			...
			...
			...

Output

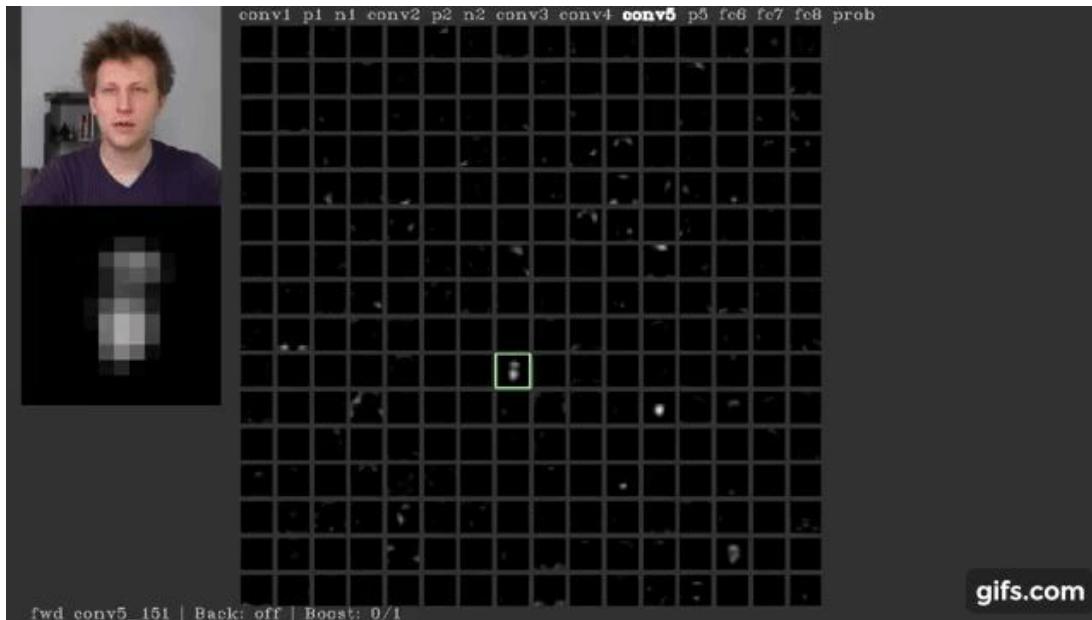


Convolutional Neural Networks (CNN)

- Pioneered by Yann LeCun back in 1989, why didn't it take off then?



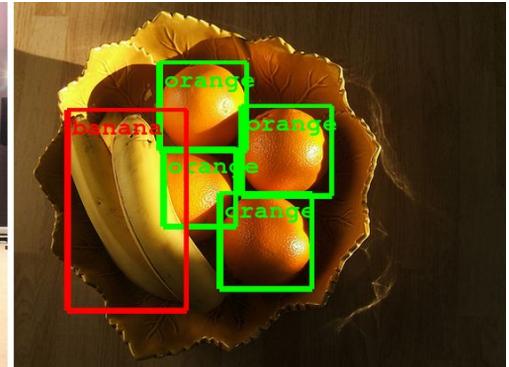
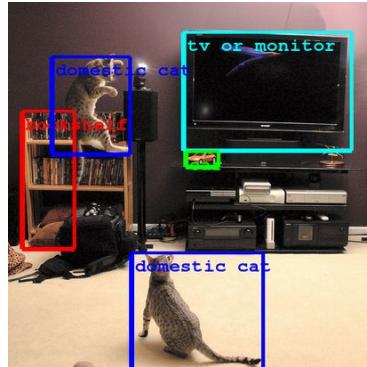
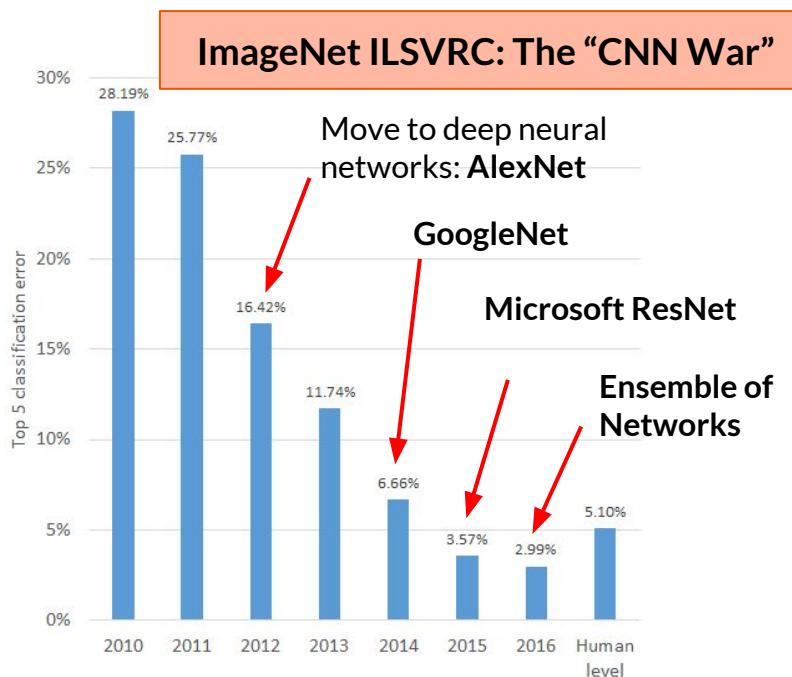
Convolutional Neural Networks (CNN)



- CNNs can be visualized to gain better understanding of what the network is learning!

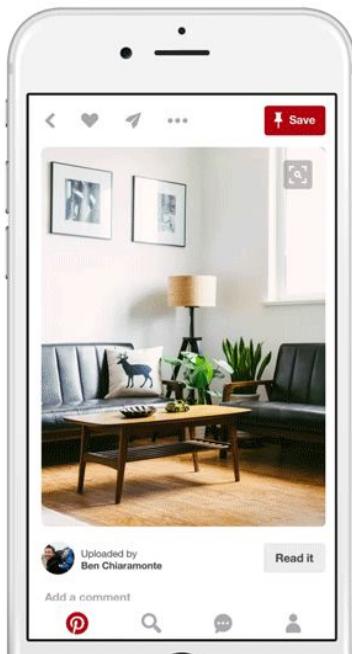
<https://github.com/yosinski/deep-visualization-toolbox>

Object Recognition and Detection/Localization

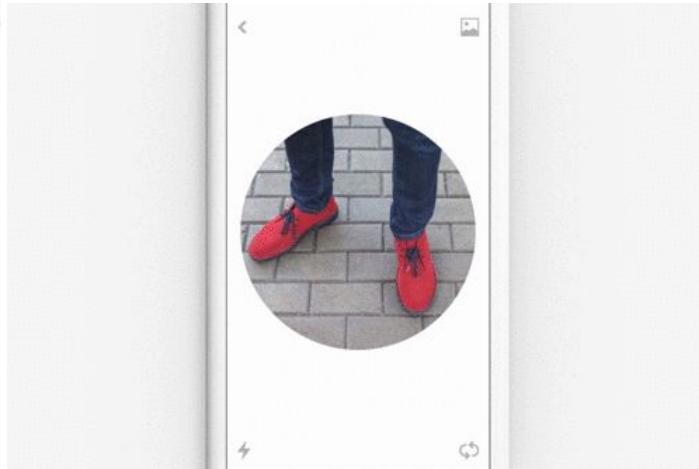


For Detecting objects:
R-CNN, Fast R-CNN, YOLO

Use Case: Visual Search @ Pinterest



Powered by CNNs!

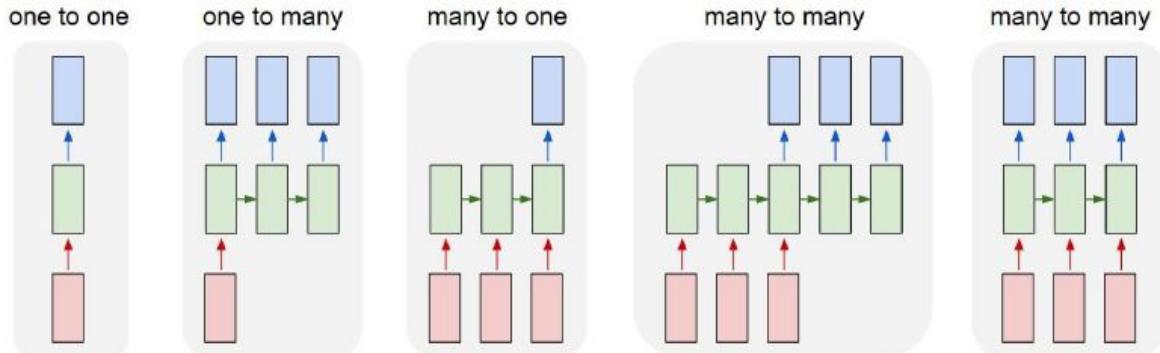


What was used?

- Object detection
- Object classification
- Recommendation engine:
“Related Pins”

Recurrent Neural Networks (RNN)

- Pioneered by Juergen Schmidhuber and Sepp Hochreiter
- Great choice if your **data changes over time**, or **patterns are sequential**
- RNNs can map input sequences to output sequences!



Recurrent Neural Networks (RNN)

- Suitable for: Time series prediction, handwriting recognition, speech recognition, stock market forecasting, text translation, sentiment analysis

Single input, Sequential output

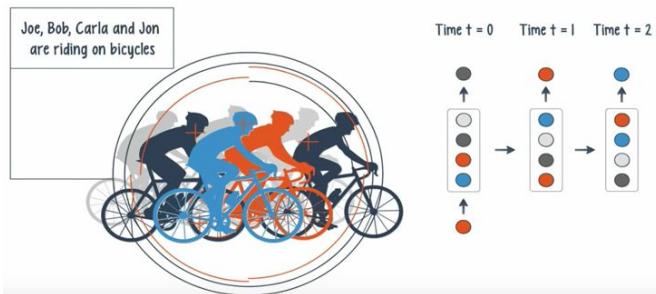


Image captioning

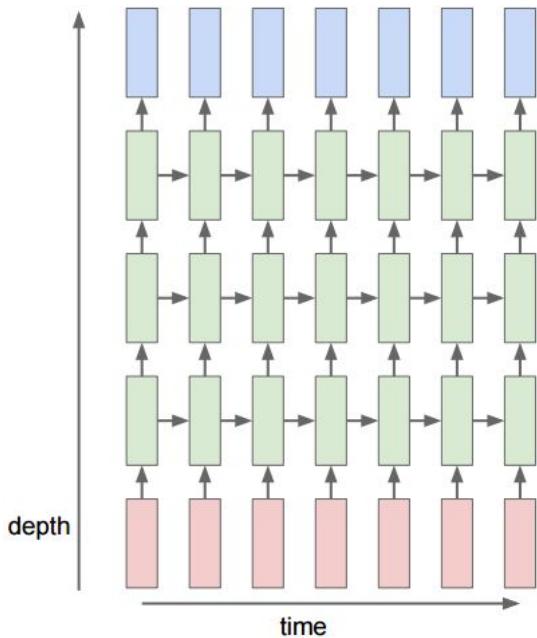
Sequential input, Single output



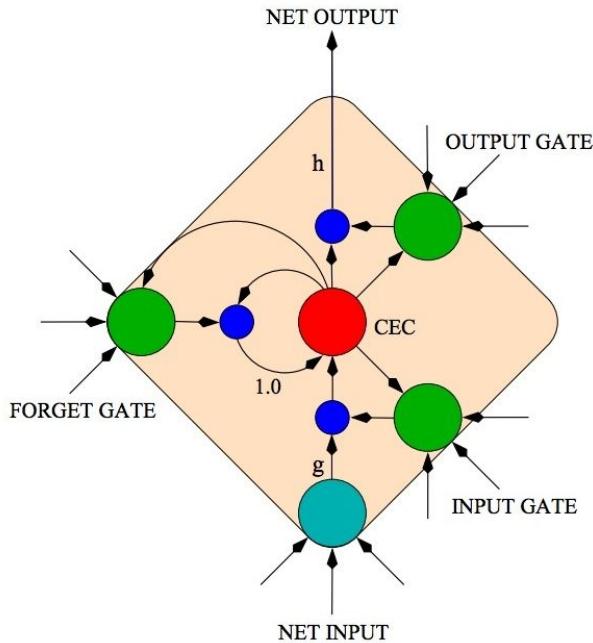
Sentiment analysis

Recurrent Neural Networks (RNN)

- Challenges:
 - Difficult to train, sensitive to parameters
 - Vanishing gradient problem just got worse!
(A 100 time step RNN = 100 layer MLP!)
- RNNs can also be stacked into multiple layers to create more complex representations



Long Short-Term Memory (LSTM) RNN



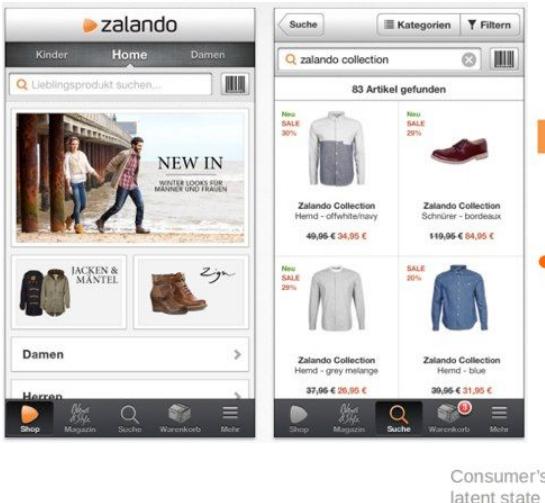
- **LSTM** is a variant of RNN that provides better back-propagation dynamics,
- LSTM units give the network **memory cells** to read, write and reset operations. During training, the network can learn when it should “remember” and when it should “forget”

LSTM: Modeling Text Data

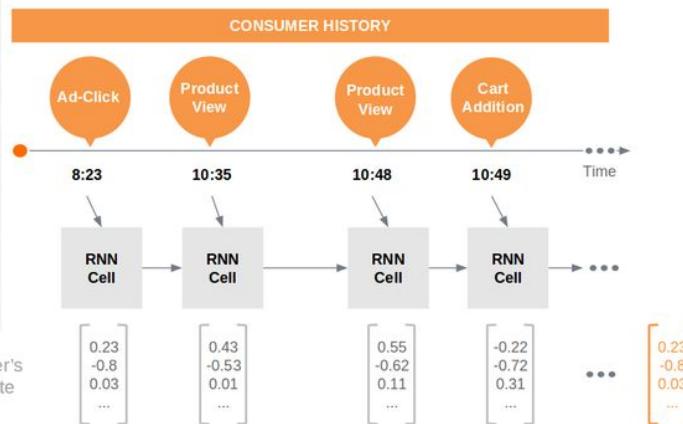
This is one of Crichton's best books. The characters of Karen Ross, Peter Elliot, Munro, and Amy are beautifully developed and their interactions are exciting, complex, and fast-paced throughout this impressive novel. And about 99.8 percent of that got lost in the film. Seriously, the screenplay AND the directing were horrendous and clearly done by people who could not fathom what was good about the novel. I can't fault the actors because frankly, they never had a chance to make this turkey live up to Crichton's original work. I know good novels, especially those with a science fiction edge, are hard to bring to the screen in a way that lives up to the original. But this may be the absolute worst disparity in quality between novel and screen adaptation ever. The book is really, really good. The movie is just dreadful.

- **Text is sequential:** Insert characters into LSTM to predict probabilities of the next letter
- Text models can be turned into **sentiment classifiers**: Cater to longer bunch of sentences, like product/book/movie reviews

Use Case: Forecasting Consumer Preferences @Zalando



Powered by RNNs!

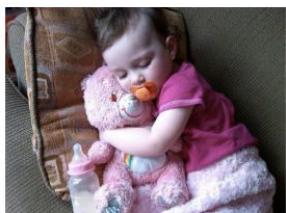
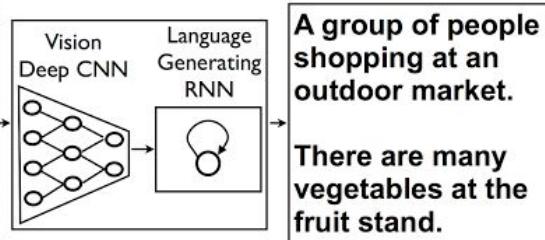


What was achieved?

Track consumer sessions

“Memorize” patterns in consumer behaviour to predict probability of orders

Hybrid Networks



A close up of a child holding a stuffed animal



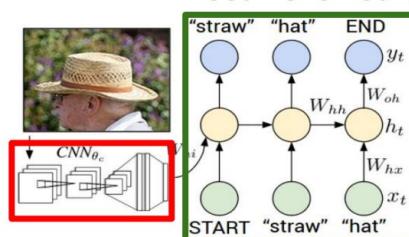
Two pizzas sitting on top of a stove top oven



A man flying through the air while riding a skateboard

- **Image Captioning:** Combines CNN and RNN
 - Generates fitting natural language captions only based on the pixel data

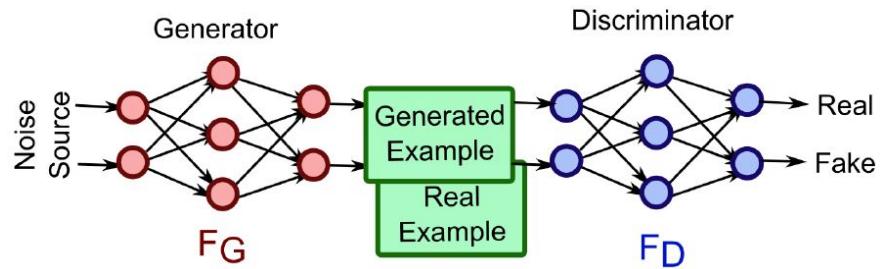
Recurrent Neural Network



Convolutional Neural Network

Generative Adversarial Networks (GAN)

- Inspired by adversarial competition in **Game Theory**
 - **Two networks**: One tasked to generate content, the other to judge content



GAN: Text-to-Image Synthesis

This flower has small, round violet petals with a dark purple center

$$\varphi \downarrow \quad \varphi(t)$$

$$z \sim N(0, 1)$$

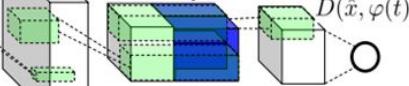
$$\hat{x} := G(z, \varphi(t))$$

Generator Network

This flower has small, round violet petals with a dark purple center

$$\varphi \downarrow$$

$$D(\hat{x}, \varphi(t))$$



Discriminator Network

a group of people on skis stand on the snow.



a table with many plates of food and drinks



a man in a wet suit riding a surfboard on a wave.



two plates of food that include beans, guacamole and rice.



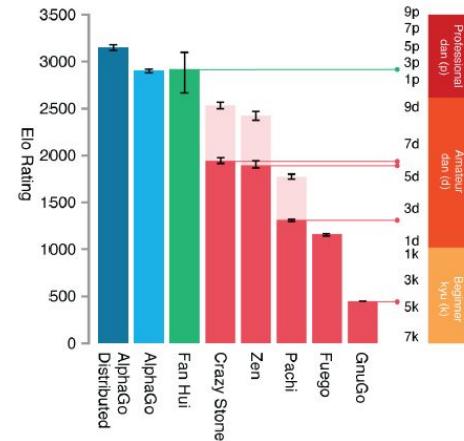
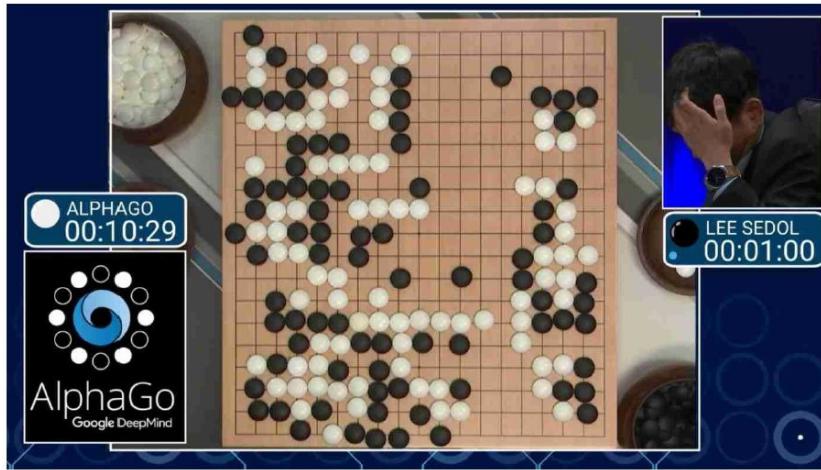
a pitcher is about to throw the ball to the batter.



a picture of a very clean living room.



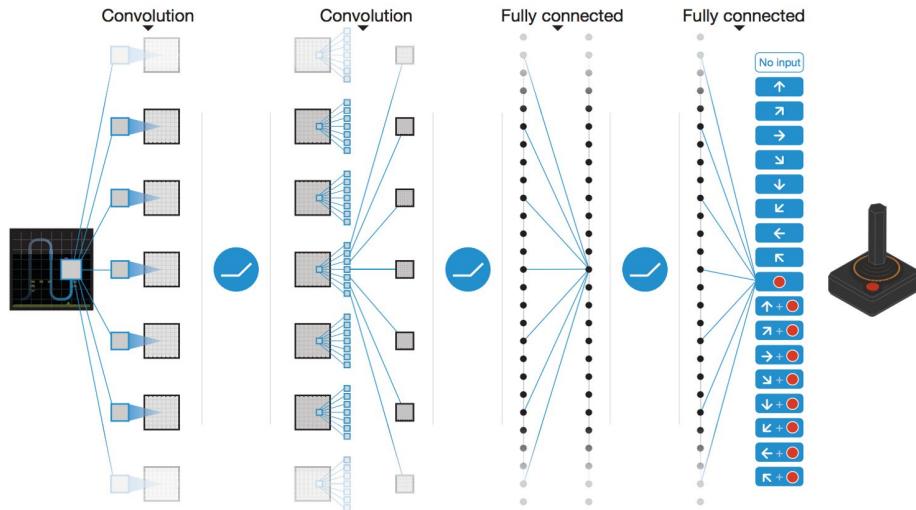
Deep Learning for Games



Examines thousands of human GO moves, and then masters it by playing with itself over and over again.

History was made: Google DeepMind's **AlphaGo** beats Go champion Lee Sedol

Deep Learning for Games



Just In: OpenAI's bot beats professional Dota 2 players



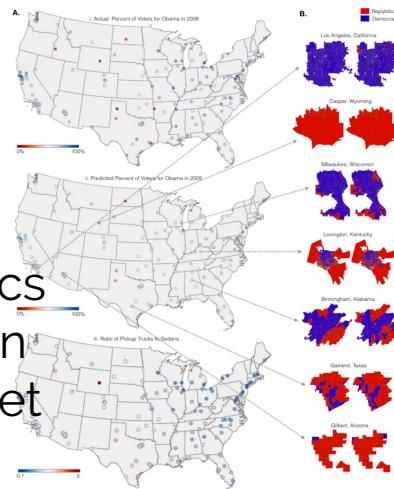
- **Deep Q-Network (DQN)** is a model-free approach to reinforcement learning using deep networks in environments with discrete action choices

What Else Can Deep Learning Generate?

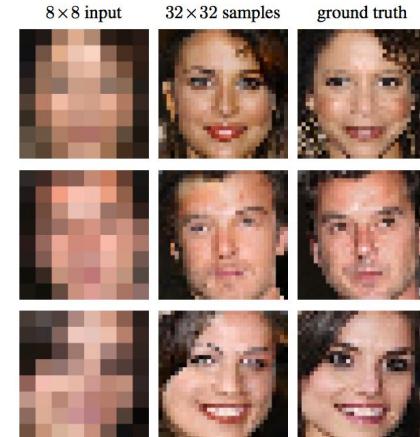


= Generate Expressive Music

Learning Demographics
and Predicting Election
Results based on Street
View Images

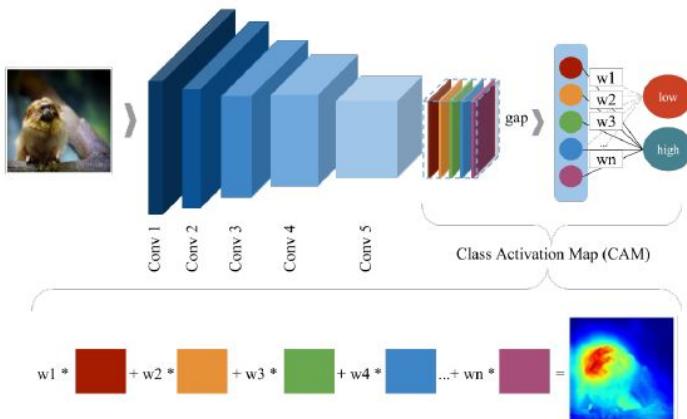


Pixel Restoration “CSI-style”

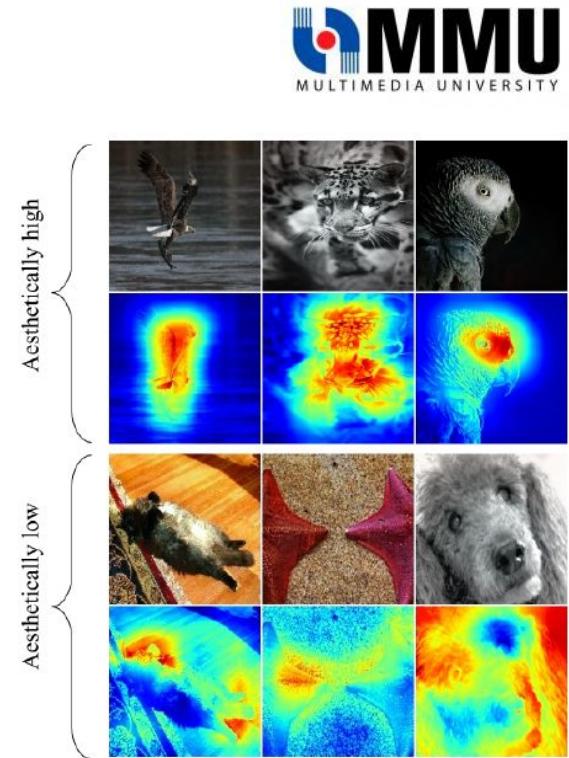


Some Homegrown Research

Predicting and Localizing Aesthetics in Images

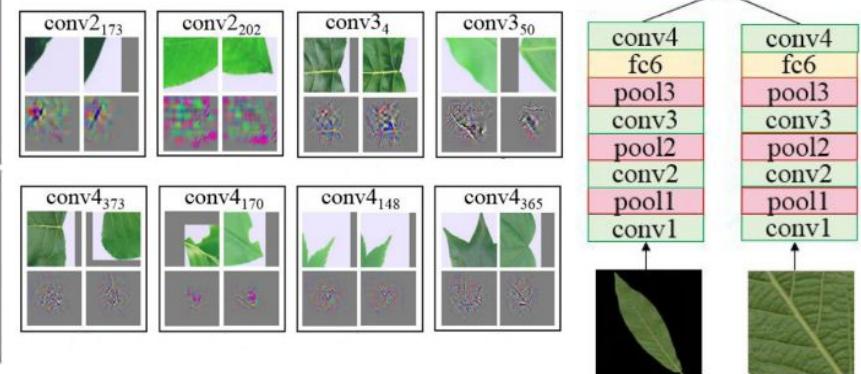
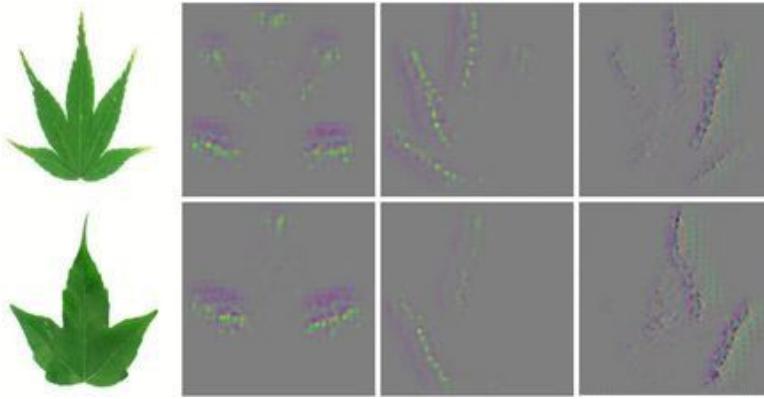


State-of-the-art:
82.27% (Text-augmented MultGAP Network, ICIP 2017)



Some Homegrown Research

Plant Identification & Learning Leaf Features



MalayaKew and Flavia datasets

Top-1 accuracy: 96.3% (Early Fusion conv-sum, PR 2017)

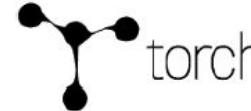
Deep Learning Tools - Open Source!



DL4J Deep Learning for Java



theano



K Keras



Caffe



Lasagne

dmlc
mxnet

NVIDIA DIGITS



Is Deep Learning over-hyped?

Is Deep Learning a black box that is unexplainable?

What should I do if I want to get started?

Demo

TensorFlow & Keras

Applied Deep Learning: Course Preview

Deep Learning: 5-day short course

- Fundamental concepts, Hands-on practicals
- Basic knowledge in Python programming & ML (added advantage)
- Tentative course outline:
 1. Machine Learning Primer - Classification & Regression
 2. Fundamentals of Neural Networks (MLPs, Shallow Networks)
 3. Deep Networks I: Convolutional Neural Networks
 4. Deep Networks II: Recurrent Neural Networks
 5. Tricks and Techniques in Deep Learning (Regularization, Data Augmentation, Ensembles)

John See
johnsee@gmail.com

Poo Kuan Hoong
kuanhoong@gmail.com



A large word cloud centered around the words "thank you" in various languages. The words are rendered in different colors and sizes, creating a dense and colorful composition. The languages represented include German (danke), Chinese (謝謝), French (merci), Spanish (gracias), English (thank you), Russian (спасибо), Portuguese (obrigado), Polish (dziękuje), Korean (감사합니다), and many others from around the world.

Evaluation

Please fill up the online evaluation form for today's preview session

<http://bit.ly/magicfeedback>

